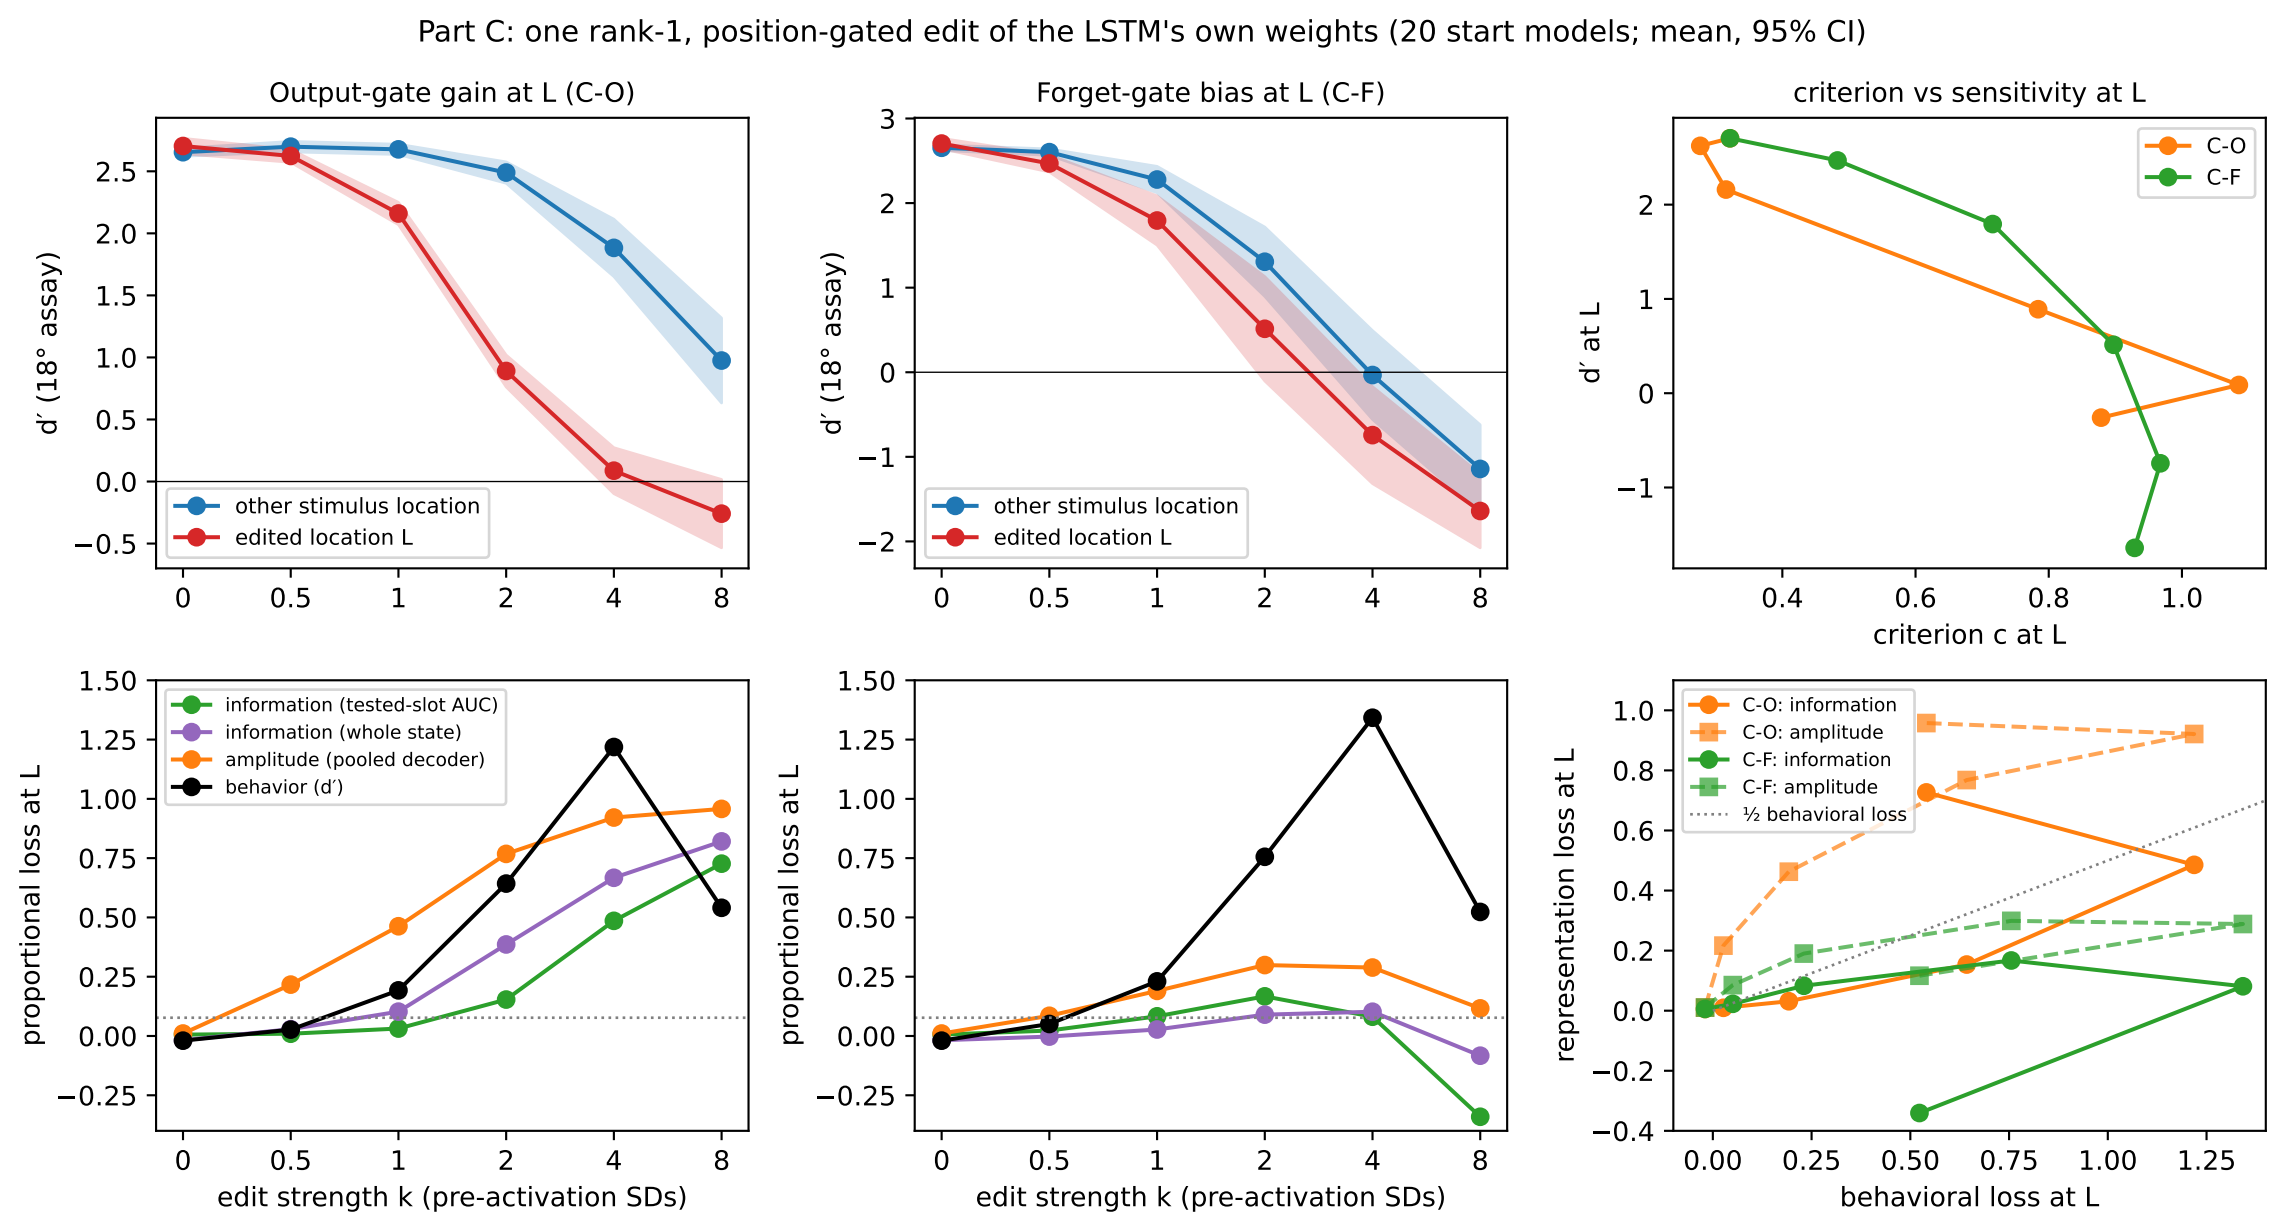
Part C: one rank-1, position-gated edit of the LSTM's own weights (20 start models; mean, 95% CI)
Output-gate gain at L (C-O)
Forget-gate bias at L (C-F)
criterion vs sensitivity at L
d′ (18° assay)
other stimulus location
edited location L
d′ at L
criterion c at L
C-O
C-F
information (tested-slot AUC)
information (whole state)
amplitude (pooled decoder)
behavior (d′)
proportional loss at L
edit strength k (pre-activation SDs)
C-O: information
C-O: amplitude
C-F: information
C-F: amplitude
½ behavioral loss
representation loss at L
behavioral loss at L