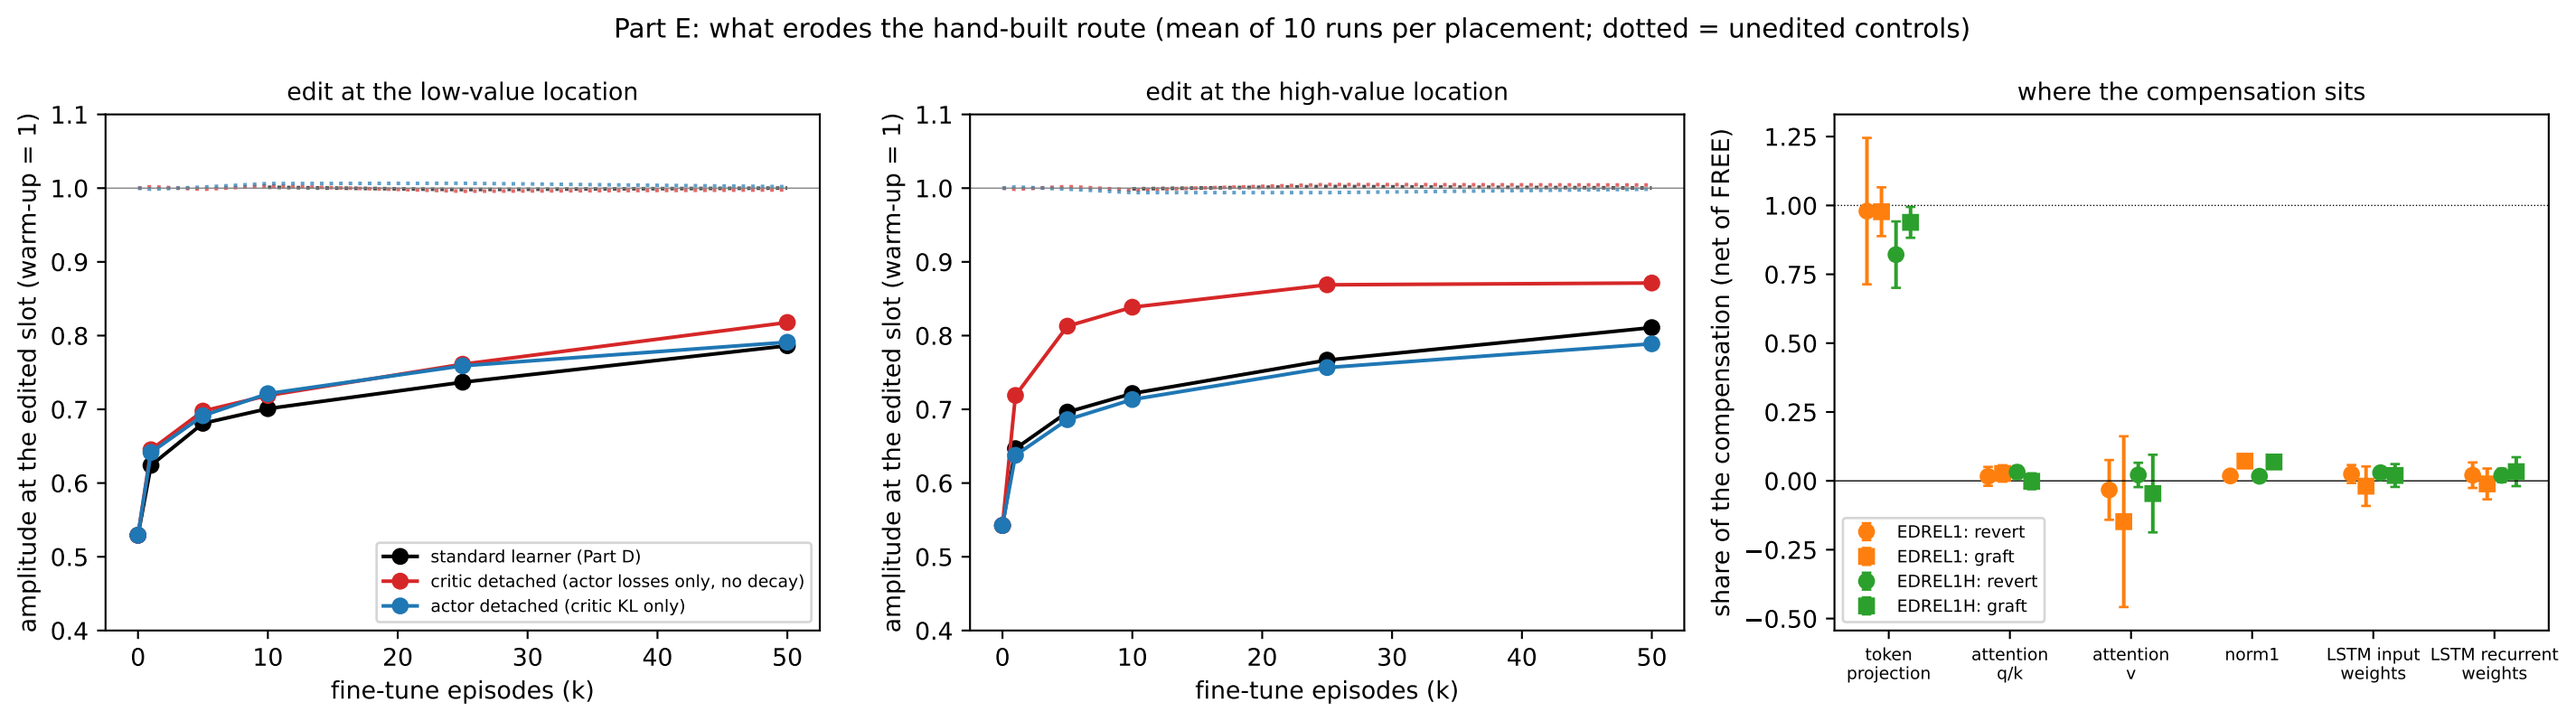
Part E: what erodes the hand-built route (mean of 10 runs per placement; dotted = unedited controls)
edit at the low-value location
edit at the high-value location
where the compensation sits
amplitude at the edited slot (warm-up = 1)
fine-tune episodes (k)
standard learner (Part D)
critic detached (actor losses only, no decay)
actor detached (critic KL only)
share of the compensation (net of FREE)
EDREL1: revert
EDREL1: graft
EDREL1H: revert
EDREL1H: graft
token projection
attention q/k
attention v
norm1
LSTM input weights
LSTM recurrent weights