


Architecture / spatial memory inside the encoder



Schematic of the implemented KDA arm. Each accumulator output is concatenated with its convolutional feature map before the next block (or final flatten). All learned parameters were trainable. This is not a spatial attention visualization.