


Downstream 7×7 recurrent readout — distinct measurements



Actual trial0 at frame14. These are ConvGRU write/reset gates and state energy, not KDA gates, implicit coefficients or attention probabilities. Channel averaging is stated explicitly. Full per-timestep arrays retained; no intermediate head is scored as an in-distribution decision.